

When Users Don't Want Privacy: Identity Laundering and Behavioral Privacy Nihilism in Generative AI

Timipado Imomotebegha
Loughborough University
Loughborough, Leicestershire, United Kingdom
t.silikowei-imomotebegha@lboro.ac.uk

Abstract

Privacy frameworks assume disclosure is a cost users seek to minimize. These assumptions fail to explain a growing phenomenon: users voluntarily broadcasting their most intimate AI conversations as social media content. This paper introduces Behavioral Privacy Nihilism, the active celebratory embrace of surveillance as identity performance, and Identity Laundering, the design mechanism by which platforms transform intimate data into shareable tokens. Through a case study of ChatGPT's "Your Year" feature in December 2025, we document how platforms institutionalize organic privacy eroding behaviors into celebratory rituals. We argue the primary harm is not data leakage but norm erosion: voluntary disclosure becoming the default, creating conditions for Identity Capture. Current governance frameworks, designed for external threats, cannot address harms users actively embrace.

CCS Concepts

• **Human-centered computing**; • **Human computer interaction (HCI)**; • **Security and privacy**; • **Social aspects of security and privacy**; • **Empirical studies in HCI**;

Keywords

privacy paradox, self-disclosure, generative AI, platform design, surveillance, identity, social media, ChatGPT

ACM Reference Format:

Timipado Imomotebegha. 2026. When Users Don't Want Privacy: Identity Laundering and Behavioral Privacy Nihilism in Generative AI. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3798964>

1 Introduction

For two decades, privacy researchers have documented a puzzling gap between what users say about privacy and what they actually do. Brown [1] first identified this "privacy paradox" in 2001: users express concern about privacy but continue sharing personal data when offered something in return. Subsequent research confirmed the pattern across contexts, with Kokolakis [2] concluding that privacy concerns are "not a valid predictor of privacy behavior." Various explanations emerged, including risk benefit calculations, cognitive biases, and bounded rationality [3], but the underlying

assumption remained constant: users want privacy, they just fail to protect it.

Meanwhile, disclosure norms have steadily escalated. Social Penetration Theory [4] holds that healthy self disclosure progresses gradually through relationship stages, yet social media disrupted this pattern as intimate information once reserved for close relationships now flows freely to networked audiences. Longitudinal research found that Facebook use "cultivates more relaxed privacy attitudes," subsequently increasing self disclosure in both online and offline contexts [5]. Each generation normalizes what the previous found transgressive.

In 2016, Spotify introduced a design template that would reshape this landscape: Wrapped. At a time when privacy debates were intensifying, Spotify reframed surveillance data as a gift. Users received colorful summaries of their listening habits, packaged for social sharing. Scholars have examined how Wrapped reframed surveillance data as celebratory self expression, blurring the line between data extraction and cultural participation [6]. One observer described Wrapped as "a pretty box full of mined data that we willingly fed into an international corporation's computers" [7]. Another argued Spotify makes Wrapped "viral and appealing" to "occlude the harms of the extractive practices" that enable its creation [8]. The template spread as YouTube, Apple Music, Duolingo, and dozens of other platforms launched their own year end recaps, training users to expect and celebrate annual summaries of their surveilled behavior.

In December 2025, this template crossed a threshold. OpenAI launched "Your Year with ChatGPT," applying the Wrapped format to generative AI conversations. Users received AI generated poems reflecting their "journey of growth," badges declaring them "The Navigator" or "Creative Debugger," and statistics tracking message counts and "Top 1%" rankings. But unlike Spotify, where the underlying data concerns music taste, ChatGPT conversations contain queries about failing marriages, career anxieties, mental health struggles, and late night spiraling. Users did not merely tolerate this synthesis of their vulnerabilities. They celebrated it, broadcast it, and converted it into status.

This paper documents this phenomenon and argues we have witnessed a shift beyond the privacy paradox. Existing constructs assume users want privacy but fail to achieve it. The privacy paradox describes attitude behavior inconsistency [1]. Privacy resignation captures defeat: users who would protect themselves if they could [9]. Privacy cynicism names coping through mistrust [10]. None capture what we observed: active celebratory embrace of intimate disclosure as identity performance.

We propose Behavioral Privacy Nihilism to name this orientation. Not giving up because you have to, but giving up because it is fun



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/2026/04

<https://doi.org/10.1145/3772363.3798964>

and social. The user is not a victim of a system they cannot change. The user is a protagonist who has co opted surveillance for social currency. Disclosure is the value, not the cost.

This paper makes four contributions. First, we introduce Behavioral Privacy Nihilism as a distinct orientation toward disclosure, extending existing constructs to capture active embrace rather than passive defeat. Second, we propose Identity Laundering as the design mechanism by which platforms transform intimate data into shareable tokens, explaining how users broadcast what they would never share directly. Third, through a case study of ChatGPT Wrapped, we document Platform Capture: how platforms observe organic privacy eroding behaviors and institutionalize them through feature design. Fourth, we identify a governance gap: current frameworks assume users want protection and cannot address harms users actively embrace.

2 BEHAVIORAL PRIVACY NIHILISM

Privacy scholarship has developed several constructs to explain the persistent gap between privacy attitudes and privacy behaviors. Each assumes, implicitly or explicitly, that users value privacy and would protect it under better conditions. We argue this assumption no longer holds universally, and that a distinct orientation has emerged that existing constructs cannot capture.

The privacy paradox [1] was the first attempt to name the gap. Users report valuing privacy but behave in ways that contradict those stated preferences. Researchers have proposed various explanations including cognitive biases, present bias discounting future risks, or inability to calculate complex tradeoffs [3]. In this framing, the user is inconsistent, perhaps irrational, but fundamentally privacy preferring.

Subsequent work moved beyond inconsistency to examine structural and psychological barriers. Draper and Turow [9] introduced privacy resignation, arguing that corporations have deliberately cultivated a sense of futility. Users believe resistance is pointless, so they surrender. The user would protect themselves if they could, but the system offers no viable path. Resignation, in this account, is learned helplessness applied to data protection. Hoffmann and colleagues [10] identified a related but distinct response in privacy cynicism: users develop mistrust as a psychological defense, assuming the worst about data practices and disengaging from protective efforts. Cynicism shields against disappointment but does not resolve the underlying desire for privacy. Hargittai and Marwick [11] documented yet another response in privacy apathy: users who simply do not care enough to act, whether from fatigue, indifference, or competing priorities. Unlike cynicism, apathy is passive. The user neither protects nor actively thinks about disclosure.

Despite their differences, these constructs share a common structure. Privacy remains a good that users want but fail to achieve. The failure may be cognitive, as in the paradox. It may be structural, as in resignation. It may be defensive, as in cynicism. It may be motivational, as in apathy. But in each case, disclosure remains a cost, a loss, something users tolerate rather than seek.

Behavioral Privacy Nihilism names a fundamentally different orientation. The nihilist does not fail to protect privacy. The nihilist does not want privacy protected. Disclosure is not a cost to be minimized but a resource to be leveraged. Surveillance outputs

become raw material for identity construction. The private self is not a sanctuary to be defended but a mine to be extracted for social value.

We use "nihilism" deliberately. In philosophical terms, nihilism denies inherent meaning or value in a domain previously assumed to possess it. Engelmann and Nissenbaum [12] recently applied this concept at the regulatory level, arguing that privacy nihilism emerges when platforms and regulators treat data categories as meaningless in an age of AI inference. We extend their insight to user behavior: the behavioral privacy nihilist denies that privacy itself retains value worth protecting. This is not defeat or apathy. It is a positive revaluation in which disclosure has become the value.

Three features distinguish Behavioral Privacy Nihilism from adjacent constructs. First, it is active rather than passive. The nihilist does not merely fail to protect data; they actively seek opportunities to transform surveillance into social currency. The DIY ChatGPT Wrapped prompts circulating before any official feature existed demonstrate users pursuing disclosure without platform prompting. Second, it is celebratory rather than resigned. The nihilist does not disclose reluctantly but with enthusiasm, creativity, and competitive energy. "Top 1%" rankings become status markers. AI generated personality badges become identity claims. The emotional valence is positive. Third, it reframes the user's relationship to platforms. In the paradox, resignation, and cynicism framings, the user is positioned as victim. The nihilist is a protagonist who has co opted surveillance infrastructure for their own ends. The platform extracts data; the user extracts identity.

This reframing does not mean nihilism is harmless. Protagonists can make choices that damage themselves and others. But understanding the orientation requires recognizing its internal logic: for the nihilist, privacy is not being violated. Privacy is being spent.

3 IDENTITY LAUNDERING

Behavioral Privacy Nihilism describes an orientation toward disclosure, but orientation alone does not explain how users come to share intimate data they would never reveal directly. Something intervenes between the nihilistic orientation and the public performance. We call this intervening mechanism Identity Laundering.

The term draws on financial money laundering, which transforms illegitimate funds into usable assets through three stages: placement, layering, and integration. Platform design accomplishes an analogous transformation for intimate data, converting disclosures users would never share directly into identity tokens they eagerly broadcast.

In the placement stage, users enter intimate, vulnerable queries into AI systems. The conversational interface mimics confidential exchange: one to one, responsive, nonjudgmental. Research confirms users disclose to chatbots at levels equivalent to human confidants [13], with some studies showing greater disclosure due to reduced fear of judgment [14]. Users share fears about relationships, confess professional failures, process grief. They do so because the interaction feels safe, not because they have calculated privacy tradeoffs. This data is dirty in the sense that users would not share it publicly. Posting "I asked ChatGPT if my marriage is failing" carries stigma. The vulnerability is too visible, the need too exposed.

In the layering stage, the platform transforms these raw queries through algorithmic abstraction. Individual prompts about relationship anxiety and professional inadequacy become aggregated into thematic summaries and archetypal labels. "The Navigator" collapses a year of confused questioning into a coherent identity narrative. An AI generated poem synthesizes vulnerable moments into aesthetic artifact. This abstraction accomplishes what Engelmann and Nissenbaum [12] call conceptual overfitting: complex human experiences forced into clean data categories. Yet users actively embrace it because the reductive archetype feels validating rather than diminishing. The Barnum effect [15] likely amplifies this acceptance, as people embrace vague personality descriptions as uniquely applicable when feedback appears personalized and emerges from perceived authority.

In the integration stage, users share the laundered outputs on social media. The transformation is complete: data too vulnerable to share directly has become identity content that feels safe, even desirable, to broadcast. A "Creative Debugger" badge implies a year of struggling with problems without exposing any specific struggle.

The critical insight is that laundered content still reveals the intimate data. The transformation is psychological, not informational. Aesthetic abstraction creates permission to disclose what was already there. Identity Laundering thus explains how Behavioral Privacy Nihilism becomes actionable. The orientation creates willingness to spend privacy as social currency. The mechanism provides the mint that transforms raw vulnerability into spendable tokens.

4 PLATFORM CAPTURE

Identity Laundering describes a design mechanism. But mechanisms do not emerge spontaneously. Platforms observe user behavior, identify patterns worth amplifying, and build infrastructure to scale them. We call this process Platform Capture: the institutionalization of organic privacy eroding behaviors into official features that legitimize and accelerate disclosure.

The genealogy of Wrapped illustrates this pattern. Spotify launched the first Wrapped campaign in December 2016, offering users colorful summaries of their annual listening habits [6]. The innovation was not data analysis but reframing. At a moment when public concern about surveillance was rising, Spotify presented tracking data as a gift. The campaign succeeded because it aligned with an existing desire: users wanted to perform their taste publicly but lacked a sanctioned format. Spotify provided the format, and users provided the free marketing. By 2021, Wrapped had become a cultural event generating over 60 million social media shares [16]. Other platforms followed as YouTube, Apple Music, Duolingo, Strava, and dozens more launched annual recaps, each training users to expect and celebrate yearly summaries of their surveilled behavior.

This training proved consequential when the template transferred to generative AI. In early December 2025, before OpenAI launched any official feature, users began creating DIY ChatGPT Wrapped experiences. Viral prompts circulated across Reddit, TikTok, and X instructing users how to generate year end summaries from their conversation histories [17]. Third party tools emerged:

Kapwing launched an unofficial generator where users could upload their entire exported chat history to receive visual recaps [18]. Power users leveraged ChatGPT's Canvas feature to build React dashboards visualizing their own data [19]. The behavior was organic. No platform prompted it. Users had been so thoroughly trained by years of Wrapped culture that they spontaneously applied the template to their most intimate data.

OpenAI observed and captured. On December 22, 2025, the company launched "Your Year with ChatGPT," an official feature incorporating the elements users had already been creating: archetypes, poems, statistics, visual artifacts [20]. But the official version added something the DIY version could not: eligibility requirements. To access Wrapped, users must have Reference Saved Memories enabled, Reference Chat History enabled, and must meet a minimum conversation activity threshold [20]. These are surveillance settings. By making them prerequisites for a desirable feature, OpenAI inverted the consent framing. Users who had previously disabled memory for privacy reasons found themselves reenabling it to qualify for their recap. Surveillance settings became feature requirements. Privacy tradeoffs became unlock conditions.

This sequence constitutes Platform Capture in its purest form. First, a platform trains users through years of Wrapped culture to expect and desire annual data summaries. Second, users spontaneously apply this expectation to a new context, creating DIY versions that demonstrate demand. Third, the platform observes the organic behavior and builds official infrastructure that requires deeper data access than the DIY version needed. Fourth, the platform frames the result as privacy forward and user controlled [20], neutralizing criticism while scaling the behavior. The user feels they have chosen to participate. The platform has merely provided a prettier version of what users already wanted. But the official version locks in surveillance settings and legitimizes disclosure at population scale.

Platform Capture does not force disclosure. It makes disclosure the path of least resistance while calling the result user choice.

5 FROM NORM EROSION TO IDENTITY CAPTURE

The preceding sections established three components: an orientation (Behavioral Privacy Nihilism), a mechanism (Identity Laundering), and a process (Platform Capture). This section addresses consequences. We argue the primary harm is not data extraction but norm erosion, and that norm erosion creates conditions for a deeper vulnerability we call Identity Capture.

Privacy norms are not fixed. Longitudinal research documents generational drift as each cohort normalizes sharing behaviors their predecessors found transgressive [5]. But the ChatGPT Wrapped phenomenon suggests something beyond gradual drift. The transition from DIY experimentation to official feature occurred within weeks, not generations. Users who in November 2025 might have hesitated to share AI conversations were by January 2026 broadcasting AI synthesized personality assessments as identity content. The norm did not erode slowly. It collapsed.

Three features characterize this collapse. First, the shift moved from experimental to expected. Early DIY prompts carried a tone of

playful hacking, but once OpenAI sanctioned the behavior, abstention required justification. Second, the shift moved from individual to collective. A single user sharing their archetype is personal expression; millions doing so simultaneously creates cultural infrastructure and shared expectations about appropriate behavior [21]. Third, the shift formalized and scaled a disclosure pattern already emerging organically. Official Wrapped did not create this impulse. It industrialized it, providing polished tokens where users had been crafting their own.

This norm erosion creates conditions for Identity Capture: a state in which users delegate self knowledge to platform infrastructure. The concept extends recent work on the "Algorithmic Self" [22], which argues that AI systems actively co produce identity through their interpretations. When users accept AI generated archetypes as authoritative accounts of who they are, they outsource a function previously reserved for introspection, close relationships, and therapeutic contexts.

Identity Capture operates through epistemic displacement. The AI's synthesis becomes the authoritative interpretation of the user's mental states. "I am The Navigator" is not a metaphor the user coined but a label the platform assigned. When users report finding their Wrapped "spookily accurate," they are not merely enjoying entertainment. They are granting epistemic authority to an algorithmic interpretation of their intimate data.

6 THE GOVERNANCE GAP

Current privacy governance operates on assumptions that Behavioral Privacy Nihilism systematically violates. Frameworks like the General Data Protection Regulation [23] and the EU Artificial Intelligence Act [24] were designed for a world where users want protection and harm comes from external actors. These frameworks offer little traction on harms users actively embrace.

The foundational assumption is that users are privacy preferring agents who would protect themselves given adequate information and control. Behavioral Privacy Nihilism inverts this relationship. The nihilist does not want protection. The nihilist wants better tools for converting surveillance into social currency. Transparency about data practices does not discourage disclosure; it may enhance the perceived value of participation.

Three governance gaps emerge. First, existing frameworks cannot regulate harms users want. Second, collective harms from individual choices fall outside regulatory reach: each user's decision is individually rational, but the aggregate effect is norm erosion that pressures future users to participate. Third, design patterns that manufacture consent evade frameworks focused on consent validity. Platform Capture does not invalidate consent; it shapes conditions under which consent feels obvious and socially required.

The gap between current governance and Behavioral Privacy Nihilism is structural, not technical. Current approaches are insufficient, and the shift toward nihilistic disclosure is accelerating faster than governance can adapt.

7 CONCLUSION

This paper has argued that privacy scholarship and governance face a challenge their foundational assumptions cannot accommodate:

users who do not want privacy protected. Through the case of ChatGPT Wrapped, we documented how Behavioral Privacy Nihilism transforms disclosure from cost to currency, how Identity Laundering converts intimate data into shareable tokens, how Platform Capture institutionalizes organic privacy erosion, and how these processes culminate in Identity Capture.

The implications extend beyond generative AI. The Wrapped paradigm has demonstrated that surveillance can be reframed as gift, that users can be trained to celebrate data extraction, and that this training transfers across contexts. Each successful application of the template normalizes the next.

We do not claim that all users are nihilists. But the growing population who actively embrace disclosure represents a phenomenon existing frameworks cannot address.

Acknowledgments

The authors used Claude (Anthropic) to assist with grammar, language, and copy editing. All conceptualization, theoretical contributions, and analysis are the authors' own. The authors reviewed and take full responsibility for all content.

References

- [1] Barry Brown. 2001. Studying the Internet Experience. HP Laboratories Technical Report HPL-2001-49. Hewlett-Packard, Palo Alto, CA.
- [2] Spyros Kokolakis. 2017. Privacy attitudes and privacy behaviour: A review of current research on the privacy paradox phenomenon. *Computers & Security* 64 (Jan. 2017), 122–134. <https://doi.org/10.1016/j.cose.2015.07.002>
- [3] Alessandro Acquisti and Jens Grossklags. 2005. Privacy and rationality in individual decision making. *IEEE Security & Privacy* 3, 1 (Jan. 2005), 26–33. <https://doi.org/10.1109/MSP.2005.22>
- [4] Irwin Altman and Dalmas A. Taylor. 1973. *Social Penetration: The Development of Interpersonal Relationships*. Holt, Rinehart and Winston, New York.
- [5] Mina Tsay-Vogel, James Shanahan, and Nancy Signorielli. 2018. Social media cultivating perceptions of privacy: A 5-year analysis of privacy attitudes and self-disclosure behaviors among Facebook users. *New Media & Society* 20, 1 (Jan. 2018), 141–161. <https://doi.org/10.1177/1461444816660731>
- [6] Taylor Annabell and Nina Vindum Rasmussen. 2025. An algorithmic event: The celebration and critique of 'Spotify Wrapped'. *AoIR Selected Papers of Internet Research* (2025). <https://doi.org/10.5210/spir.v2024i0.14038>
- [7] Meredith Clark. 2020. Spotify 2020 Wrapped is a trite humblebrag. But maybe it's the least bad way to end this year. *NBC News Think* (Dec. 6, 2020). Retrieved January 16, 2026 from <https://www.nbcnews.com/think/opinion/spotify-2020-wrapped-trite-humblebrag-maybe-it-s-least-bad-ncna1250148>
- [8] Rachel Metz. 2021. Spotify Wrapped shows how our personal data gets sliced and diced. *CNN Business* (Dec. 2, 2021). Retrieved January 16, 2026 from <https://edition.cnn.com/2021/12/02/tech/spotify-wrapped-data>
- [9] Nora A. Draper and Joseph Turow. 2019. The corporate cultivation of digital resignation. *New Media & Society* 21, 8 (Aug. 2019), 1824–1839. <https://doi.org/10.1177/1461444819833331>
- [10] Christian Pieter Hoffmann, Christoph Lutz, and Giulia Ranzini. 2016. Privacy cynicism: A new approach to the privacy paradox. *Cyberpsychology: Journal of Psychosocial Research on Cyberspace* 10, 4 (Dec. 2016), article 7. <https://doi.org/10.5817/CP2016-4-7>
- [11] Eszter Hargittai and Alice Marwick. 2016. "What Can I Really Do?" Explaining the Privacy Paradox with Online Apathy. *International Journal of Communication* 10 (2016), 3737–3757. <http://ijoc.org/index.php/ijoc/article/view/4655>
- [12] Severin Engelmann and Helen Nissenbaum. 2025. Countering Privacy Nihilism. In *Conceptions of Data Protection and Privacy: Legal and Philosophical Perspectives*, Elisa Orrù and Ralf Poscher (Eds.). Hart Publishing.
- [13] Johan Hoorn, Elly A. Konijn, Karolina Karpinski, and Paul Ketelaar. 2024. Self-disclosure to conversational AI: a literature review, emergent framework, and directions for future research. *Personal and Ubiquitous Computing* (Aug. 2024). <https://doi.org/10.1007/s00779-024-01823-7>
- [14] Annie Ho, Jeffrey T. Hancock, and Adam S. Miner. 2018. Psychological, Relational, and Emotional Effects of Self-Disclosure After Conversations With a Chatbot. *Journal of Communication* 68, 4 (Aug. 2018), 712–733. <https://doi.org/10.1093/joc/jqy026>
- [15] Bertram R. Forer. 1949. The fallacy of personal validation: A classroom demonstration of gullibility. *Journal of Abnormal and Social Psychology* 44, 1 (1949),

- 118–123. <https://doi.org/10.1037/h0059240>
- [16] The Dartmouth. 2024. Trends: Spotify's evolution into a social platform. Retrieved January 16, 2026 from <https://www.thedartmouth.com/article/2024/01/trends-spotifys-evolution-into-a-social-platform>
- [17] TechRadar. 2025. This prompt creates your very own ChatGPT Wrapped 2025 — here's how to get your AI year in review. Retrieved January 16, 2026 from <https://www.techradar.com/computing/artificial-intelligence/chatgpt-users-are-making-a-spotify-wrapped-for-their-year-in-ai-heres-how-to-do-it>
- [18] Kapwing. 2025. ChatGPT Wrapped 2025. Retrieved January 16, 2026 from <https://www.kapwing.com/chatgpt-wrapped>
- [19] Alex Mann. 2025. I created my own 2025 ChatGPT Wrapped (and you can too). How-To Geek (Dec. 10, 2025). Retrieved January 16, 2026 from <https://www.howtogeek.com/create-your-2025-chatgpt-wrapped/>
- [20] OpenAI. 2025. Your Year with ChatGPT - FAQs. OpenAI Help Center. Retrieved January 16, 2026 from <https://help.openai.com/en/articles/20001042-your-year-with-chatgpt-faqs>
- [21] Peyton Young. 2015. The evolution of social norms. *Annual Review of Economics* 7, 1 (2015), 359–387. <https://doi.org/10.1146/annurev-economics-080614-115322>
- [22] Jeena Joseph. 2025. The algorithmic self: how AI is reshaping human identity, introspection, and agency. *Frontiers in Psychology* 16 (June 2025). <https://doi.org/10.3389/fpsyg.2025.1645795>
- [23] European Parliament and Council of the European Union. 2016. Regulation (EU) 2016/679 of the European Parliament and of the Council on the protection of natural persons with regard to the processing of personal data and on the free movement of such data (General Data Protection Regulation). *Official Journal of the European Union* L119 (May 2016), 1–88. Retrieved January 16, 2026 from <https://eur-lex.europa.eu/eli/reg/2016/679/oj>
- [24] European Commission. 2024. Regulation (EU) 2024/1689 of the European Parliament and of the Council on artificial intelligence (AI Act). *Official Journal of the European Union* (July 2024). Retrieved January 16, 2026 from [https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=\\$CELEX:32024R1689](https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=$CELEX:32024R1689)